

Daniil Dorokhov

PRODUCT MANAGER (AI/ML)

Armenia (residence permit) · open to relocation · dorokhov.daniil69@gmail.com · t.me/ananasDDA · ananasdda.github.io/Portfolio · github.com/ananasDDA

SUMMARY

An engineer moving into AI/ML product management. I take AI out of the demo and into production: at R7 (current role: QA Automation Lead) I identified the product's bottleneck myself, then built and still own an internal platform where large language models do real work - generating tests, triaging failures, analysing logs and screenshots - with validated output, cost control, and graceful degradation when a model fails. Alongside that, as a co-founder, I shipped an AI-powered app to the App Store and launched an iOS app for on-device LLM inference. Technical foundation: four years of engineering, a business informatics degree, and computer science coursework at HSE. I am looking for a role that requires deciding where AI creates value and where it only adds cost.

EXPERIENCE

QA Automation Lead; initiator and owner of the internal test platform

02/2025 - present

R7, Corporate Server

Joined as the first engineer on the product. Formal title is QA Automation Lead; the automation platform is something I initiated, designed and run as an internal product - finding the bottlenecks, shaping the plan, and defending it to management.

- **Found the main bottleneck and removed it.** Preparing environments before release testing took three days, and the release simply waited. I brought it down to 30 minutes by automating deployment across 11 servers and 5 operating system families. Business outcome: testing starts the day the build ships, and the release cycle is three days shorter.
- **Built an internal product for test result analysis:** REST API and web interface, a pipeline running from data ingest to verdict - metric enrichment, LLM analysis of run statistics, vision-based comparison of reference and actual screenshots, a quality gate with configurable thresholds and a pass/fail verdict, webhook notifications, and a self-contained HTML report. It is wired into CI and used by three testing teams and the development team.
- **Decided where AI belongs, not merely how to adopt it.** LLMs are used where they earn their cost - test generation and failure triage - and skipped entirely when a run is green. Spend is controlled through rate limiting and a request queue, secrets are masked before any call, a model outage never fails the pipeline, and a path to on-premise inference is designed in.
- **Made the outcome visible to the B2B customer:** an HTML test plan report, executive summaries and Telegram digests are used during release acceptance. Before that, testing status existed only inside the team.
- **Proved the platform generalises:** another company product's test suite runs on the same infrastructure with no rewrite.
- **Process metrics I owned:** time-to-test three days → 30 minutes; test authoring 15-30 minutes → minutes; failure triage hours → seconds; commit feedback in about 10 minutes.
- Team: one automation engineer, one intern, and two manual testers part-time; hiring, mentoring, and department training.

QA Automation Engineer

03/2024 - 02/2025

Avroid, startup

- Spotted a problem in the support team, proposed a solution and drove it to rollout: a Python Telegram bot integrated with the CRM and issue tracker. Still in use today.
- A second product for the same team: RAG over the Aurora OS documentation plus an agent, giving support an answer grounded in the documentation and a reproduction of the scenario on a mobile Linux emulator.
- As the first QA in the company, established the testing process from scratch and took part in hiring and training support staff.

QA Engineer, release acceptance

08/2022 - 03/2024

MyOffice

- Release acceptance of B2B products: the go/no-go decision on a build, working with development, DevOps, information security and customers - stakeholder communication on both sides of a release.
- Covered the vacant DevOps role for six months.

OWN PRODUCTS

P.R.O. - AI-powered sports app

2025 - present

Co-founder, backend and AI

- Took the product from idea to publication on the App Store, including Apple review.
- I own the AI functionality: personalised recommendations, goal setting matched to fitness level, progress forecasting, and an assistant. Within the founding team I take part in deciding which AI features the user actually needs.

PRO Local AI Lab

iOS app for on-device LLM and VLM inference

- Product hypothesis: privacy and no subscription as the value proposition. It forces explicit trade-offs between model quality, speed and device memory.

SKILLS

Product: bottleneck discovery and prioritisation, process metrics (time-to-test, cycle time, release readiness), requirements and acceptance, stakeholder and B2B customer communication, App Store launch

AI in product: assessing where language and multimodal models (LLM, VLM) apply, guardrails and output validation, cost control, degradation without service failure, on-device inference and its trade-offs, cloud versus on-premise

Technical foundation: Python, Java, Kotlin, REST API, Kafka, CI/CD, Docker, Kubernetes, AWS, Linux - I assess feasibility and speak the same language as engineering

Tools: TestIT, Jira, Confluence, Git, Cursor and AI assistants in daily work

Languages: Russian - native; English - Upper-Intermediate (B2)

EDUCATION

BSc, Business Informatics

Plekhanov Russian University of Economics

A degree at the intersection of business process and IT: process modelling, requirements analysis, business economics.

Auditing student

HSE University, Faculty of Computer Science

Calculus, discrete mathematics, algorithms and data structures, computer architecture and operating systems; machine learning and recommender systems in progress.